

## Git Repo : [https://github.com/ashokitschool/kubernetes\\_manifest\\_yaml\\_files.git](https://github.com/ashokitschool/kubernetes_manifest_yaml_files.git)

=====  
Kubernetes (K8S)  
=====

=> It is free & open source s/w.

=> Google developed this k8s using "GO" programming language.

=> K8S provides Orchestration (Management) platform.

=> K8S is used to manage containers  
(create/stop/start/delete/scale-up/scale-down)

=====  
Advantages  
=====

1) Container Orchestration : Manages containers.

2) Self Healing : If any pod got crashed/deleted then it will replace with new pod.

3) Load Balancing : Load will be distributed all containers which are up and running.

4) Auto Scaling : Based on demand containers count will be increased or decreased.

=====  
Docker Vs Kubernetes  
=====

Docker : Containerization platform

Note: Packaging our app code and dependencies as single unit for execution is called as Containerization.

Kubernetes : Orchestration platform

Note: Orchestation means managing the containers.

=====  
Kubernetes Architecture  
=====

=> K8S will follow cluster architecture.

=> Cluster means group of servers will be available.

=> In K8s Cluster we will have master node (control plane) and worker nodes.

=====  
K8S Cluster Components  
=====

1) Control Node (Master Node)

- API Server
- Scheduler
- Controller Manager
- ETCD

## 2) Worker Nodes

- Kubelet
- Kube Proxy
- Docker Runtime
- POD
- Container

=> To deploy our application using k8s then we need to communicate with control node.

=> We will use KUBECTL (CLI) to communicate with control plane.

=> "API Server" will receive the request given by "kubectl" and it will store that request in "ETCD" with pending status.

=> "ETCD" is an internal database of k8s cluster.

=> "Scheduler" will identify pending requests available in ETCD and it will identify worker node to schedule the task.

Note: Scheduler will identify worker node by using kubelet.

=> "Kubelet" is called as Node Agent. It will maintain all the worker node information.

=> "Kube Proxy" will provide network for the cluster communication.

=> "Controller Manager" will verify all the tasks are working as expected or not.

=> In Worker Node, "Docker Engine" will be available to run docker container.

=> In K8s, container will be created inside POD.

=> POD is a smallest building block that we can create in k8s cluster to run our applications.

Note: In K8s, everything will be represented as POD only.

=> Pods are used to deploy applications and manage their lifecycle within a Kubernetes environment.

=====  
K8S Cluster Setup  
=====

=> We can setup k8s cluster in multiple ways

### 1) Self Managed k8s cluster (We need to setup everything)

- a) Mini Kube => Single Node Cluster => Only for beginners practice
- b) Kubeadm => Multi Node Cluster (HA)

### 2) Provider Managed K8S cluster (ready made)

- a) AWS EKS
- b) Azure AKS
- c) GCP GKE

Note: Provider will give ready made cluster for us.

#### Note: Provider Managed Clusters are chargeable ####

=====  
Assignment : Setup MiniKube software  
=====



```
id: 68686
name: Smith
address:
  city: Hyd
  state: TG
  country: India
```

```
=====
K8S Manifest YML
=====
```

=> To deploy our application in kubernetes we need MANIFEST YML file.

=> In k8s manifest YML we will write 4 sections.

```
apiVersion: <resource-version-number>
```

```
kind: <resource-type>
```

```
metadata: <resource-info>
```

```
spec: <container-info>
```

```
=====
POD Maniest YML
=====
```

```
---
apiVersion: v1
kind: Pod
metadata:
  name: javawebapppod
  labels:
    app: javawebapp
spec:
  containers:
    - name: javac1
      image: ashokit/javawebapp
      ports:
        - containerPort: 8080
...
```

Note: Save the above content in .yaml file.

```
$ kubectl get pods
```

```
$ kubectl apply -f <yaml-file-name>
```

```
$ kubectl get pods
```

```
$ kubectl logs <pod-name>
```

```
# display in which worker node our pod is running
$ kubectl get pods -o wide
```

Note: By default PODS can be accessed only with in the cluster. We can't access PODS outside.

=> To access PODS outside the cluster we need to expose the PODS by using K8S Services concept.

```
=====
K8S Service
=====
```

=> K8S service is used to group the pods and expose them for outside access.

=> In K8S we have 3 types of services

- 1) Cluster IP
- 2) Node PORT
- 3) Load Balancer

-----  
=> When we deploy our application in k8s then PODS will be created for our application..

=> For every POD one IP will be generated.

=> PODS are short lived objects. We should not access PODS using POD IP because PODS can be destroyed and re-created at any point of time.

### CLUSTER IP : It generates one static IP for all PODS based on POD label.

Ex: 192.168.10.89

Note: using ClusterIP we can access PODS only within the cluster.

Note: PODS created , PODs Destroyed, PODS increased or decreased still ClusterIP will not change.

Ex: Database PODS we should access only within in the cluster. Outside ppl should not access DB pods. In this scenario we can use CLUSTER IP service.

## Node PORT : It is used to expose the PODS for outside access.

Using NODE PUBLIC IP we can access PODS running in that node outside of the cluster also.

## LOAD BALANCER : It is used to expose the PODS for outside access.

=> When we can access LOAD BALANCER URL, it will distribute the load to all the nodes and all the PODS available for our application.

-----  
=====

K8S Service Manifest YML

=====

```
---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: NodePort
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30070
...
```

Note: Save this data with .yaml extension

```
$ kubectl get svc
```

```
$ kubectl apply -f <svc-manifest-yml>
```

```
$ kubectl get svc
```

Note: Enable NODE PORT in security group inbound rules.

Note: To access our application use below URL

URL : <http://worker-node-public-ip:node-port/java-web-app/>

=====  
What is NodePort ?  
=====

=> When we use service type as NodePort then k8s will use one random port number to expose our application on worker node for public access.

Node Port Range : 30,000 - 32767

Note: If we can also fix nodeport number in service manifest yml.

=====  
POD and Service in Single Manifest YML  
=====

```
---
apiVersion: v1
kind: Pod
metadata:
  name: javawebapppod
  labels:
    app: javawebapp
spec:
  containers:
    - name: javac1
      image: ashokit/javawebapp
      ports:
        - containerPort: 8080
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
spec:
  type: NodePort
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30070
```

...

\$ kubectl apply -f <manifest-yml-name>

\$ kubectl get pods

\$ kubectl get svc

=====  
K8S namespaces  
=====

=> Namespaces are used to group the resources logically.

mysql-db-pods =====> mysql-db-ns

backend-app-pods =====> backend-ns

frontend-app-pods =====> frontend-ns

=> Inside k8s cluster we can create multiple namespaces.

=> Each namespace is isolated with another namespace.

Note: When we delete a namespace all the resources belongs to that namespace also gets deleted.

```
# display k8s namespaces
kubectl get ns
```

```
# get pods of specific namespace
kubectl get pods -n kube-system
```

Note: In kubectl command if we don't specify any namespace then it will consider "default" namespace.

=> In K8s we can create namespace in 2 ways

- 1) using "kubectl create ns" command
- 2) Using manifest YML file

```
-----
Approach-1 :
-----
```

```
# create namespace
kubectl create ns ashokitns
```

```
# delete namespace
kubectl delete ns ashokitns
```

Note: When we delete a namespace all the resources belongs to that namespace also gets deleted.

```
-----
Approach-2 :
-----
```

```
---
apiVersion: v1
kind: Namespace
metadata:
  name: ashokit-backend-ns
...
```

```
=====
Namespace + POD + Service
=====
```

```
---
apiVersion: v1
kind: Namespace
metadata:
  name: ashokit
---
apiVersion: v1
kind: Pod
metadata:
  name: javawebapppod
  namespace: ashokit
  labels:
    app: javawebapp
spec:
  containers:
```

```

- name: javac1
  image: ashokit/javawebapp
  ports:
    - containerPort: 8080
---
apiVersion: v1
kind: Service
metadata:
  name: javawebappsvc
  namespace: ashokit
spec:
  type: NodePort
  selector:
    app: javawebapp
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30070
...

```

```

# run above yml file
kubectl apply -f <yml>

# check pods
kubectl get pods -n ashokit

# check service
kubectl get svc -n ashokit

# check all resources
kubectl get all -n ashokit

```

=====  
=> When we create POD directly using "kind: Pod" in manifest yml then k8s will not manage our pod life cycle.

=> If we delete pod then k8s will not create new POD in this scenario. Self Healing will not work in this way.

=> If we want k8s to manage POD life cycle then we should use k8s resources to create the PODS.

- 1) ReplicationController (RC) (outdated)
- 2) ReplicaSet (RS)
- 3) Deployment
- 4) DaemonSet
- 5) StatefulSet

=====  
What is ReplicationController  
=====

=> It is one of the resource in k8s to manage pod life cycle.

Note: If any pod is deleted/crashed/damaged then RC will perform self healing.

=> Using RC we can perform PODS count scale up and scale down.

```

---
apiVersion: v1

```



```

kind: ReplicationController
metadata:
  name: javawebrc
spec:
  replicas: 2
  selector:
    app: javawebapp
  template:
    metadata:
      labels:
        app: javawebapp
    spec:
      containers:
        - name: javawebct
          image: ashokit/javawebapp
          ports:
            - containerPort: 8080
...

```

```
$ kubectl apply -f rc.yml
```

```
$ kubectl get pods
```

```
$ kubectl delete pod <pod-name>
```

```
$ kubectl get pods
```

```
$ kubectl scale rc javawebrc --replicas=5
```

```
$ kubectl scale rc javawebrc --replicas=1
```

```

=====
ReplicaSet
=====

```

=> It is one of the resource in k8s to manage pod life cycle.

Note: If any pod is deleted/crashed/damaged then RS will perform self healing.

=> Using RS we can perform PODS count scale up and scale down.

```

---
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: javawebrs
spec:
  replicas: 2
  selector:
    matchLabels:
      app: javawebapp
  template:
    metadata:
      name: javawebapppod
      labels:
        app: javawebapp
    spec:
      containers:
        - name: javawebappcontainer
          image: ashokit/javawebapp
          ports:
            - containerPort: 8080
...

```

```
kubectl apply -f rc.yml
```

```
kubectl get all
```

```
kubectl get pods
```

```
kubectl delete pod <pod-name>
```

```
kubectl get pods
```

```
kubectl scale rs javawebrs --replicas=5
```

```
kubectl scale rs javawebrs --replicas=1
```

```
=====
Deployment
=====
```

=> It is one of the resource in k8s to manage pod life cycle.

=> This is the most recommended approach to deploy our applications in k8s cluster.

=> In Deployment we have 2 strategies to create PODS

1) Rolling Update

2) ReCreate

=> ReCreate means it will delete all existing pods and will create new pods. In This ReCreate approach, application will have downtime.

=> RollingUpdate means it will delete old pod and create new pod one after other. Here no downtime for our application.

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: javawebdeploy
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
  selector:
    matchLabels:
      app: javawebapp
  template:
    metadata:
      name: javawebapppod
    labels:
      app: javawebapp
    spec:
      containers:
        - name: javawebappcontainer
          image: ashokit/javawebapp
          ports:
            - containerPort: 8080
...

```

```
kubectl apply -f rc.yml
```

```
kubectl get all
```

```
kubectl get pods
```

```
kubectl delete pod <pod-name>

kubectl get pods

kubectl scale deployment javawebdeploy --replicas=5

kubectl scale deployment javawebdeploy --replicas=1
```

```
=====
HPA (Horizontal POD Scalar)
=====
```

=> HPA is used to scale up and scale down our pods automatically based on traffic.

=> HPA will not work directly, we need to configure HPA in k8s cluster.

@@@ Reference Video : <https://www.youtube.com/watch?v=c-tsJrcB50I>

```
=====
What is DaemonSet ?
=====
```

=> It is one of the resource in k8s which is used to manage PODS life cycle.

=> Automatically adds the pod to new nodes as they join the cluster.

=> Removes the pod from nodes when they are removed.

=> Ensures one pod per node (unless explicitly configured otherwise).

#### Use Cases: Running cluster-wide background services like:

=> Log collection agents (e.g., Fluentd, Filebeat, LogStash)

=> Monitoring agents (e.g., Prometheus Node Exporter)

=> Metrics Server

```
# Create fluentd pods using daemonset
$ kubectl apply -f https://k8s.io/examples/controllers/daemonset.yaml

# check fluentd pods
$ kubectl get pods -n kube-system

# delete fluentd pods by deleting daemonset we have created
$ kubectl delete daemonset fluentd-elasticsearch -n kube-system
```

```
=====
what is statefulset
=====
```

=> It is one of the resource in k8s which is used to manage PODS life cycle.

=> StatefulSet is used to manage stateful application related pods.

stateless pod = when POD deleted POD data also gets deleted.

statefull pod = when pod delete POD data will not be deleted.

Note: To make our POD as stateful we need to use Storage for the POD.

Key Features of StatefulSet:

## Stable Pod Names : Pods are created with persistent names like myapp-0, myapp-1, etc. Even if a pod is deleted, its name remains the same when it is recreated.

## Stable Storage : Each pod gets a persistent volume that is tied to it. The volume remains even if the pod is deleted and re-created.

Use Cases:

- 1) Databases (MySQL, PostgreSQL, MongoDB)
- 2) Distributed systems (Kafka, Cassandra, Elasticsearch)

=====

Blue - Green Deployment Approach

=====

=> Blue green deployment is an application release model to the production.

=> It reduces risk and minimizes downtime

=> It uses two production environments, known as Blue and Green

=> The old version can be called the blue environment (v1)

=> The new version can be known as the green environment (v2)

=====

Advantages

=====

=> Rapid releasing

=> Simple rollbacks

=> Seamless customer experience

=> Zero Downtime

-----

## Step-1 : Create blue deployment (pods will be created with label as v1)

## Step-2 : Check pods status

## Step-3 : Create Live Service to expose blue pods (Type : load balancer)

## Step-4 : Use Load Balancer URL To access our app which is running in blue pods.

URL : http://lbr-url/java-web-app/

## Step-5 : Create Green Deployment (pods will be created with latest docker image and label as v2).

## Step-6 : Verify green pods status

## Step-7 : Make green pods as live by changing "live-service" selector as 'v2' in yml file. After changing yml then re-execute live-service yml file.

## Step-8 : Use Load Balancer URL To access our app which is running in green pods.

URL : <http://lbr-url/java-web-app/>

=====  
ConfigMap & Secrets  
=====

=> For every application multiple environments will be available.

- 1) DEV
- 2) SIT
- 3) UAT
- 4) PILOT
- 5) PROD

=> Every env will have its own config properties to run the application.

Ex :

- 1) database props
- 2) smtp props
- 3) kafka server properties
- 4) redis server properties
- 5) payment gateways
- 6) third party apis urls

=> If we configure above properties within the application then our application will become tightly coupled.

=> When we want to deploy our application in another environment then we have to "change properties + re-package + re-create docker image + re-deployment". It is a time taking process and risky.

=> If we want to deploy our app in multiple environments then we need to make sure our application is loosely coupled with env properties.

Note: We need to keep environment properties outside of the application code.

## To make our app loosely coupled with env properties we can use ## ConfigMap & Secrets ##

=> Using configmap and secret we can de-couple application code and application properties so that our docker images will become loosely coupled.

Note: At the time of deployment, we can supply environment properties to the application container using ConfigMap & Secrets.

=> ConfigMap & Secret will store data in key-value format

=> ConfigMap is used to store non-sensitive data.

=> Secret is used to store sensitive data.

```
=====
ConfigMap manifest YML
=====
```

```
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: ashokit-cg-dev
data:
  db_url: "jdbc:mysql://localhost:3306/"
  db_name: "ashokit"
  db_port: "3306"
...
```

```
$ kubectl get configmap
$ kubectl apply -f <yaml>
$ kubectl get configmap
```

```
=====
Secret manifest YML
=====
```

```
---
apiVersion: v1
kind: Secret
metadata:
  name: ashokit-secret-dev
type: Opaque
data:
  db_username: YXNob2tpdA== #root
  db_password: YWJjQDEyMw== #abc@123
...
```

```
$ kubectl get secret
$ kubectl apply -f <yaml>
$ kubectl get secret
```

```
=====
how to read the data from configmap and secret
=====
```

```
env:
- name: MYSQL_DATABASE
  valueFrom:
    configMapKeyRef:
      name: ashokit-cg-uat
      key: db_name
```

```
env:
- name: MYSQL_PASSWORD
  valueFrom:
    secretMapKeyRef:
      name: ashokit-secret-dev
      key: db_password
```

```
=====
Assignment
=====
```

1) Deploy MySQL database in k8s cluster by using configmap and secret.

DB username : read from config map  
DB pwd : read from secret

1) create configmap (ex: mysql-configmap.yml)

```
---
apiVersion: v1
kind: ConfigMap
metadata:
  name: mysql-config-map
data:
  MYSQL_DATABASE: mydatabase
  MYSQL_USER: myuser
...
```

2) create secret (ex: mysql-secret.yml)

```
---
apiVersion: v1
kind: Secret
metadata:
  name: mysql-secret
type: Opaque
data:
  MYSQL_ROOT_PASSWORD: cm9vdA==
  MYSQL_PASSWORD: cm9vdA==
...
```

3) create deployment (ex: mysql-deployment.yml)

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mysql-deployment
spec:
  replicas: 1
  strategy:
    type: Recreate
  selector:
    matchLabels:
      app: mysql
  template:
    metadata:
      labels:
        app: mysql
    spec:
      containers:
        - name: mysql-container
          image: mysql:latest
          ports:
            - containerPort: 3306

      env:
        - name: MYSQL_DATABASE
          valueFrom:
            configMapKeyRef:
              name: mysql-config-map
              key: MYSQL_DATABASE

        - name: MYSQL_USER
          valueFrom:
            configMapKeyRef:
```

```
name: mysql-config-map
key: MYSQL_USER
```

```
- name: MYSQL_ROOT_PASSWORD
  valueFrom:
    secretKeyRef:
      name: mysql-secret
      key: MYSQL_ROOT_PASSWORD
```

```
- name: MYSQL_PASSWORD
  valueFrom:
    secretKeyRef:
      name: mysql-secret
      key: MYSQL_PASSWORD
```

```
...
```

```
# To get pods
```

```
$ kubectl get pods
```

```
# Go enter into POD
```

```
$ kubectl exec -it <pod-name> -- /bin/bash
```

```
# Connect with mysql database
```

```
$ mysql -u root -proot
```

```
# display databases available
```

```
$ show databases;
```

```
# select database
```

```
$ use mydatabase;
```

```
# show tables
```

```
$ show tables;
```

```
# exit from the mysql
```

```
$ exit
```

```
$ exit
```

```
=====
```

```
What is HELM ?
```

```
=====
```

```
=> In linux we will use package managers to install a software
```

```
Ex : yum , apt, rpm etc...
```

```
=> HELM is a package manager which is used to install required softwares in k8s cluster.
```

```
Ex: Metrics Server, Promethues, Grafana, ELK stack....
```

```
=> HELM will use charts to install required packages.
```

```
=> Chart means collection of configuration files (manifest ymls).
```

```
=====
```

```
Helm Installation
```

```
=====
```

```
curl -fsSl -o get_helm.sh https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3
```

```
chmod 700 get_helm.sh
```

```
./get_helm.sh
```

```
helm version
```

```
=====
```



## Kubernetes Monitoring

=====

=> We can monitor our k8s cluster and cluster components using below softwares

1) Prometheus

2) Grafana

=====

### Prometheus

=====

-> Prometheus is an open-source systems monitoring and alerting toolkit.

-> Prometheus collects and stores its metrics as time series data

-> It provides out-of-the-box monitoring capabilities for the k8s orchestration platform.

=====

### Grafana

=====

-> Grafana is an analysis and monitoring tool.

-> It provides visualization for monitoring.

-> It provides charts, graphs, and alerts for the web when connected to supported data sources.

Note: Grafana will connect with Prometheus for data source.

## Note: Using HELM charts we can easily deploy Prometheus and Grafana in K8S Cluster

=====

### Install Prometheus & Grafana In K8S Cluster using HELM

=====

# Add the latest helm repository in Kubernetes

\$ helm repo add stable https://charts.helm.sh/stable

# Add prometheus repo to helm

\$ helm repo add prometheus-community https://prometheus-community.github.io/helm-charts

# Update Helm Repo

\$ helm repo update

# install prometheus & grafana

\$ helm install stable prometheus-community/kube-prometheus-stack

# Get all pods

\$ kubectl get pods

Note: You should see prometheus pods running

# By default prometheus and grafana services are available within the cluster as ClusterIP, to access them outside lets change it to LoadBalancer by editing service directly.

\$ kubectl edit svc <service-name>

# Edit Prometheus Service & change service type to LoadBalancer then save and close that file

\$ kubectl edit svc stable-kube-prometheus-sta-prometheus

=> Access Prometheus server using below URL

URL : <http://LBR-DNS:9090/>

# Now edit the grafana service & change service type to LoadBalancer then save and close that file

```
$ kubectl edit svc stable-grafana
```

=> Access Grafana server using below URL

URL : <http://LBR-DNS/>

=> Use below credentials to login into grafana server

UserName: admin

Password: prom-operator

=====

Assignment : EFK stack setup in k8s cluster to monitor application logs

@@ Reference video : [https://youtu.be/8MLcbbfEL1U?si=bQ\\_BrOv3EiLu48eu](https://youtu.be/8MLcbbfEL1U?si=bQ_BrOv3EiLu48eu)

=====

=====

Node Selector

=====

=> Node Selector is used to schedule the pods on particular worker node only.

=> To achieve this we can assign label for the worker node and we will configure that node label in our manifest yaml as node-selector.

```
# configure label for worker node
$ kubectl get nodes
$ kubectl edit node <node-name>
```

```
# configure below label under labels section
name: ashokit-wn-1
```

Note-1: If node-selector is matching with worker-node label then our pods will be created on that particular worker-node only.

Note-2: If node-selector not matching with worker-node label then pods will not be scheduled for execution.

=> Execute below manifest yaml to create nginx deployment with 3 pod replicas

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
```

```

template:
  metadata:
    labels:
      app: nginx
  spec:
    nodeSelector:
      name: ashokit-wn-1
    containers:
      - name: nginx
        image: nginx:1.14.2
        ports:
          - containerPort: 80
...

```

```
$ kubectl apply -f <yaml>
```

```
$ kubectl get pods -o wide
```

```

=====
Node Affinity
=====

```

=> Node Affinity preferred approach.

=> If nodeSelector is matching with any worker-node label then schedule pods on that worker node only.

=> If matching is not found then schedule pods on any available worker-node in the cluster.

```

---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      affinity:
        nodeAffinity:
          preferredDuringSchedulingIgnoredDuringExecution:
            - weight: 1
              preference:
                matchExpressions:
                  - key: name
                    operator: In
                    values:
                      - ashokit
                      - ait
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
...

```

```
$ kubectl apply -f <yaml>
```

```
$ kubectl get pods -o wide
```

```
=====
Taints
=====
```

=> Taints are used to make worker node not eligible for pods scheduling and pods execution.

=> We have 3 popular taint options

1) No Schedule : Kubernetes will not schedule new pods on the node.

2) No Execute : New PODS will not be scheduled and Existing pods also will be removed.

3) Prefer No schedule : Kubernetes tries to avoid scheduling pods on the node but will schedule them if no better options exist.

```
# create taint on worker node
```

```
$ kubectl taint nodes <node-name> color=blue:NoSchedule
```

```
# remove taint
```

```
$ kubectl taint nodes <node-name> color-
```

```
=====
Tolerations
=====
```

=> Tolerations are used to schedule pods on tainted worker nodes also.

Taint: "I don't want pods on me unless they tolerate me."

Toleration: "I'm okay to run on a node with this taint."

```
# tainting worker node
```

```
$ kubectl taint nodes <node-name> key1=value1:NoSchedule
```

Note: If any pod is having tolerations as key1 and value1 then schedule those pods even though the nodes are tainted with NoSchedule state.

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      tolerations:
        - key: "key1"
          operator: "Equal"
          value: "value1"
          effect: "NoSchedule"
      containers:
        - name: nginx
```

```
image: nginx:1.14.2
ports:
- containerPort: 80
```

```
...
```

```
=====
Liveness & Readiness Probes
=====
```

Probes are checks performed by the Kubernetes kubelet to monitor the health of your application running inside a pod.

#### ## Liveness Probe :

Purpose: Indicates whether the pod is alive or dead.

If it fails: Kubernetes kills the container and may restart it based on the restart policy.

Used to detect: Crashes, deadlocks, or stuck apps.

#### ## Readiness Probe :

Purpose: Indicates whether the pod is ready to serve traffic.

If it fails: The pod is removed from Service endpoints, so it won't receive requests.

Used during: App startup and while running.

```
=====
```

- 1) What is Orchestration and Why
- 2) K8S introduction
- 3) K8S Advantages
- 4) Kubernetes Architecture
- 5) K8S Setup (AWS EKS)
- 6) PODS
- 7) Services (ClusterIP, NodePort & LoadBalancer)
- 8) Namespaces
- 9) ReplicationController
- 10) ReplicaSet
- 11) Deployment
- 12) HPA
- 13) Blue Green Deployment
- 14) ConfigMap & Secret
- 15) MySQL Deployment in k8s cluster
- 16) HELM Charts
- 17) Prometheus & Grafana

- 18) EFK Stack
- 19) What is Node Selector
- 20) What is Node Affinity
- 21) What are Taints & Tolerations
- 22) Liveness & Readiness Probes